

.NET 4.5 Task Parallel Library Dataflow Networks

Alon Fliess
CodeValue



Agenda

- Introduction to Dataflow Networks
- “Around the Block”
- Building the Network
- Concurrency Control
- Advanced TDF Topics

Task Parallel Library

- Task & Task Scheduler are the building blocks of the .NET parallel execution engine
 - You can use task to spawn parallel execution, or
 - You can use higher level abstract such as
 - PLINQ & Data Parallelism (Parallel.For, ...)
 - Task Parallel Library Dataflow Network ➔ TDF
- TDF
 - A new way to abstract sequences of code execution over tasks

The Actor Model

- "The Actor model adopts the philosophy that *everything is an actor*. This is similar to the *everything is an object* philosophy used by some object-oriented programming languages, but differs in that object-oriented software is typically executed sequentially, while the *Actor model is inherently concurrent*."

Wikipedia

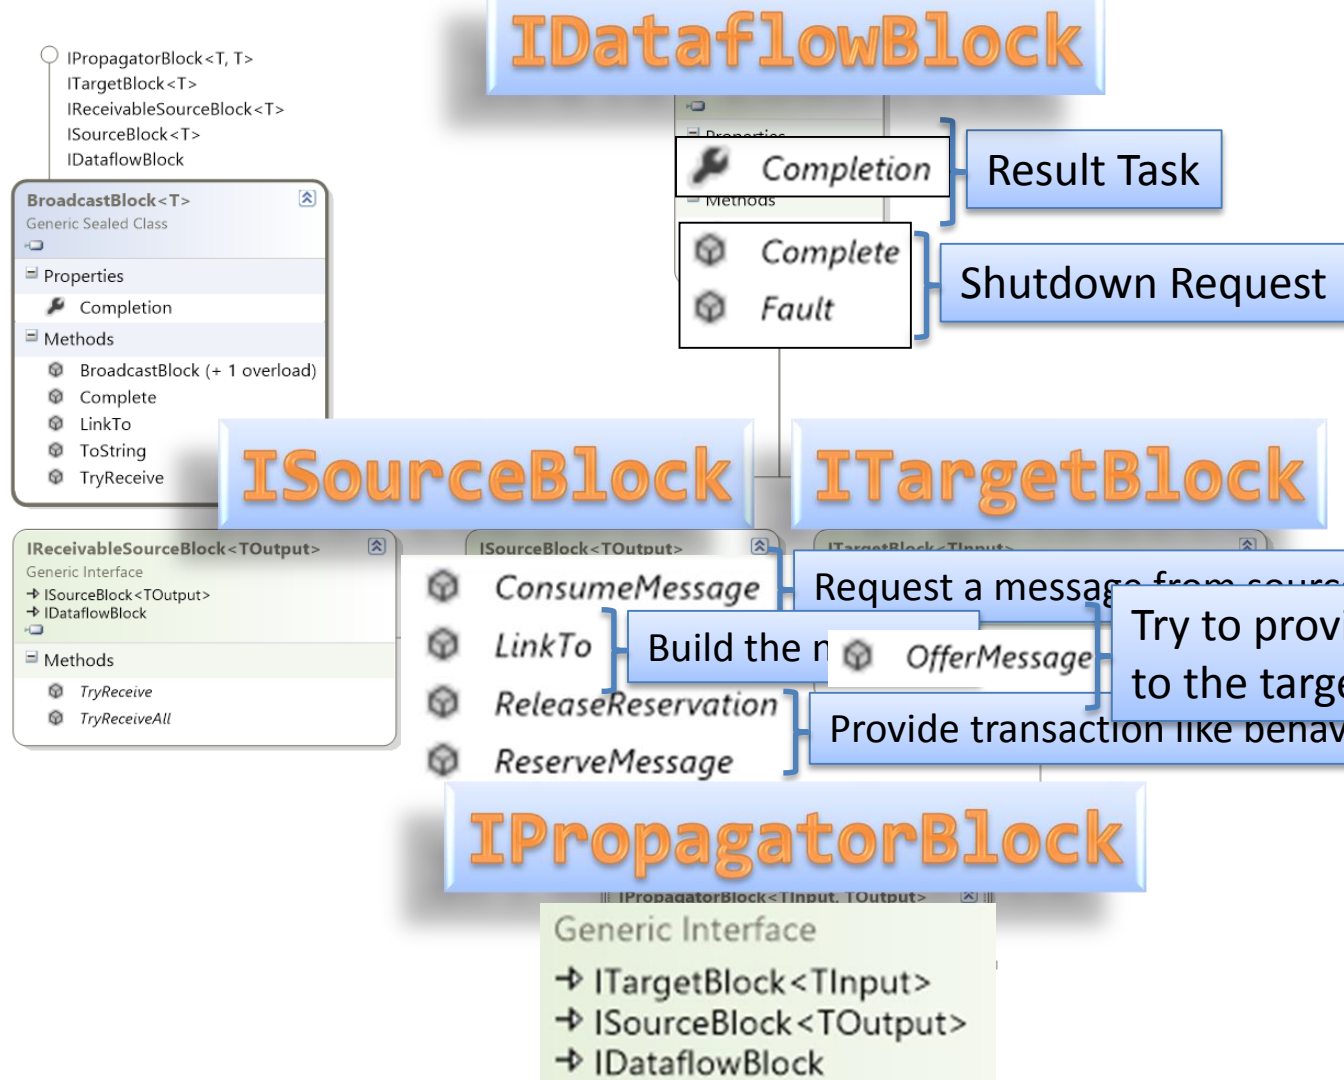
Introduction to Dataflow Networks

- TDF is a flavor of the actor model
 - It is the managed implementation of the VC++ Asynchronous Agents Library
- TDF is used to create and execute a Dataflow network
 - A good example of such a network is Microsoft Office Excel
 - Any pipeline is a simple Dataflow network
- TDF is a technology for message passing and parallel execution of CPU and I/O intensive applications
- TDF promises high performance, high throughput with low latency

“Around the Block”

- TDF network is built from message passing dataflow blocks
 - [IDataflowBlock](#)
- There are three types of blocks:
 - Source blocks
 - [ISourceBlock<out TOutput>](#)
 - [IReceivableSourceBlock<TOutput>](#)
 - Propagator Blocks
 - [IPropagatorBlock<in TInput, out TOutput>](#)
 - Target Blocks
 - [ITargetBlock<in TInput>](#)



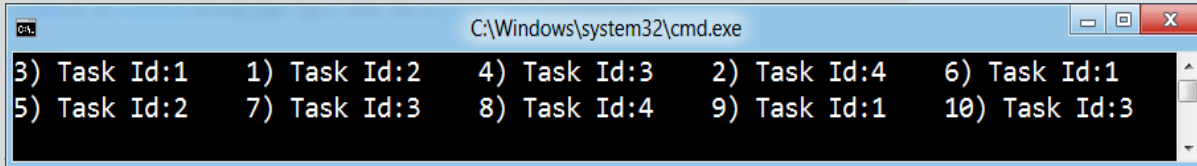


ActionBlock<in Tin>

- This is a simple target block
 - It receives data and acts upon it

```
var ab = new ActionBlock<int>(i =>
{
    Console.WriteLine("{0}) Task Id:{1}\t", i, Task.CurrentId);
    Thread.Sleep(10);
},
new ExecutionDataflowBlockOptions
{ MaxDegreeOfParallelism = 4 });

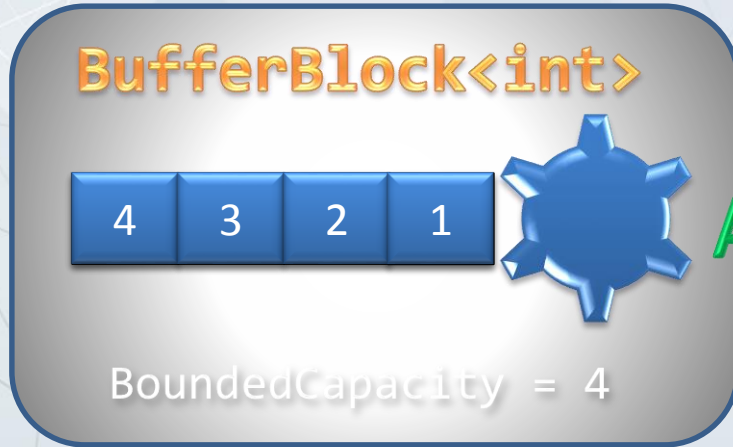
for (int i = 1; i <= 10; i++)
{
    ab.Post(i);
}
ab.Complete();
ab.Completion.Wait(),
```



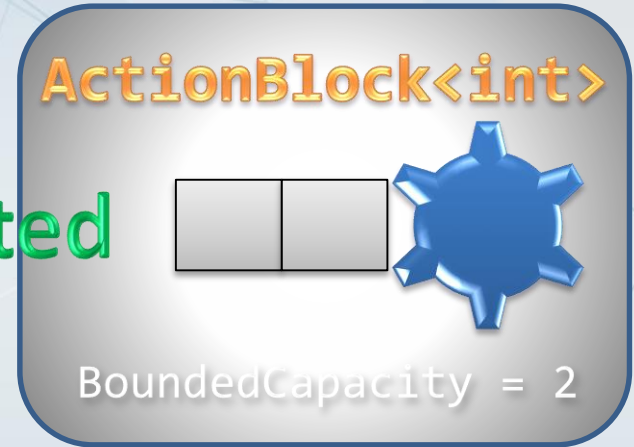
```
C:\Windows\system32\cmd.exe
3) Task Id:1    1) Task Id:2    4) Task Id:3    2) Task Id:4    6) Task Id:1
5) Task Id:2    7) Task Id:3    8) Task Id:4    9) Task Id:1    10) Task Id:3
```


Message Passing - Handshaking

OfferMessage

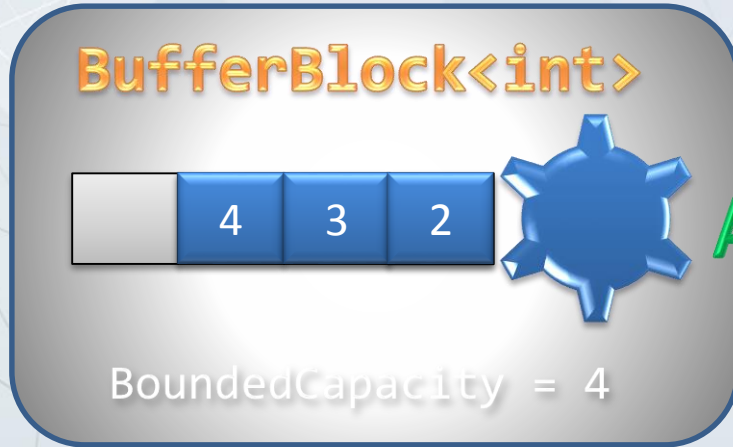


Accepted

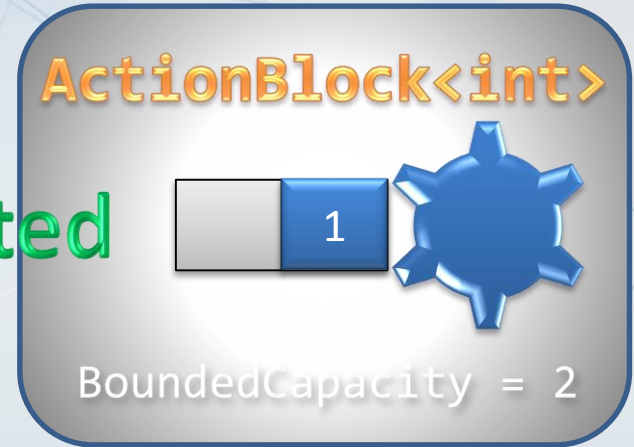


Message Passing - Handshaking

OfferMessage

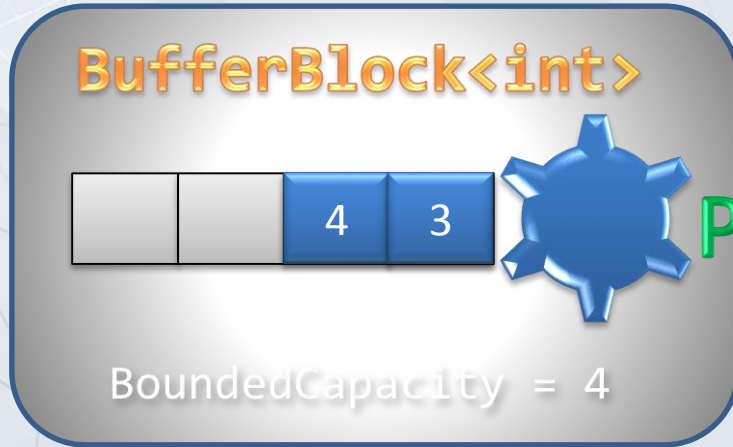


Accepted

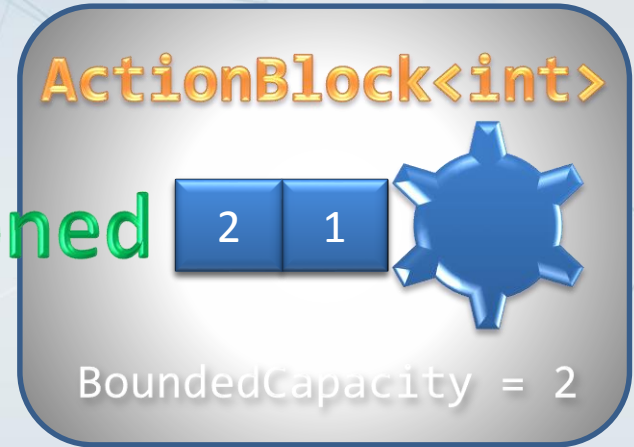


Message Passing - Handshaking

OfferMessage

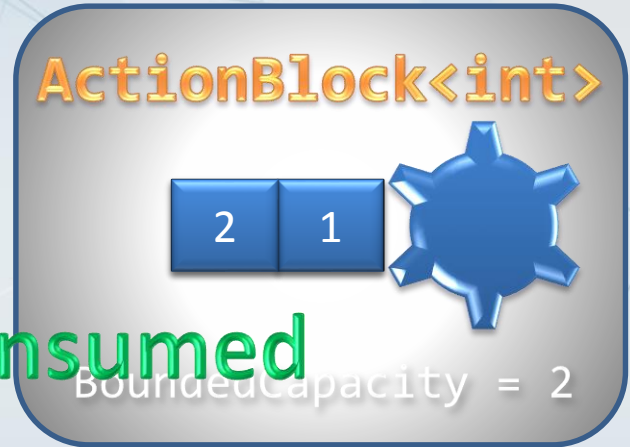
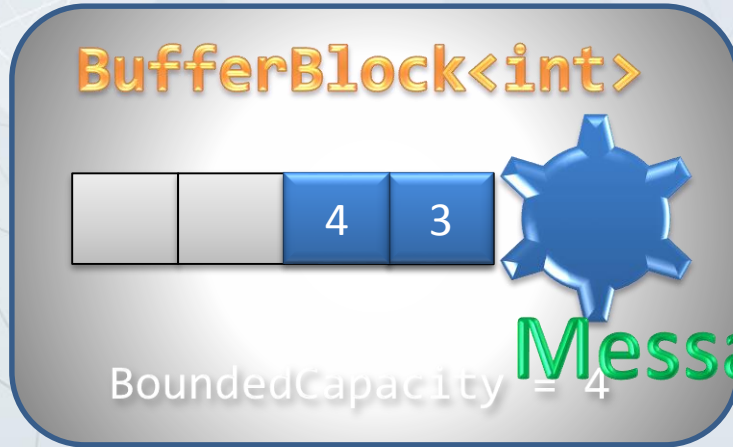


Postponed



Message Passing - Handshaking

ConsumeMessage



MessageConsumed

Message Passing - Handshaking

OfferMessage

`ab.Post(5);`

5

`ISourceBlock = null`

Declined

`ActionBlock<int>`

3

2



`BoundedCapacity = 2`

Message Passing - Handshaking

OfferMessage

ab.Post(6);

6

Declining Permanently

ISourceBlock = null

ActionBlock<int>

3

2



BoundedCompletedly = 2

ab.Complete();



Building the Network

- Building the network is a simple sequence call to:
 - `ISourceBlock<T0>.LinkTo(ITargetBlock<TOutput>, DataflowLinkOptions)`
- Or to one of the [DataflowBlock](#) extension methods:
 - `LinkTo<T0>(ISourceBlock<T0>, ITargetBlock<T0>)`
 - `LinkTo<T0>(ISourceBlock<T0>, ITargetBlock<T0>, Predicate<TOutput>)`
 - `LinkTo<TOutput>(ISourceBlock<TOutput>, ITargetBlock<TOutput>, DataflowLinkOptions, Predicate<TOutput>)`

Dataflow Link Options

Name	Description
<u>Append</u> (Links Order)	Determines whether the link should be appended or prepended to the source's list of links
<u>MaxMessages</u> (Link Life Time)	Determines the maximum number of messages that may be consumed across the link
<u>PropagateCompletion</u> (Teardown behavior)	Gets or sets whether the linked target will have completion and faulting notification automatically propagated

Link Predicate

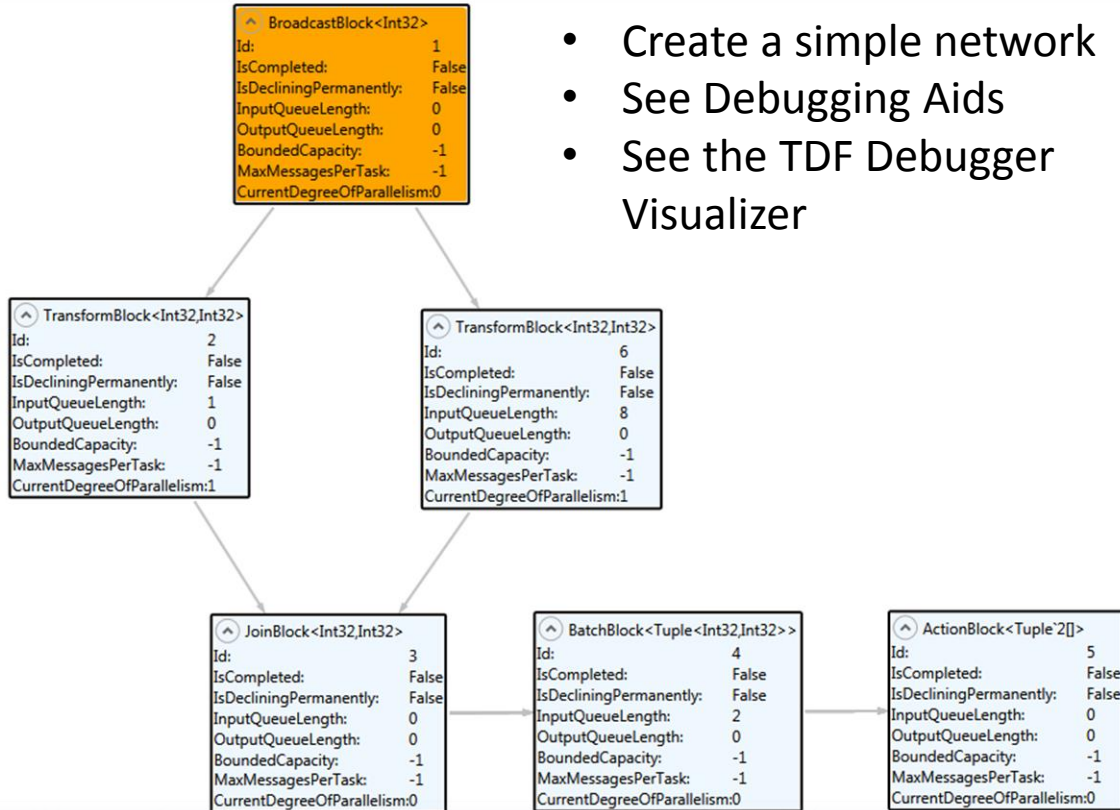
- You can supply a predicate to the [LinkTo](#) extension method

```
src.LinkTo(target, n => n % 2 == 0);
```

- Internally [LinkTo](#) creates new private **FilteredLinkPropagator** block and inserts it as a filter just before the target block
- The **FilteredLinkPropagator** block executes the predicate for all offered message
 - If the predicate returns true, it offers the message to the target block
 - Otherwise it declines the message

Simple Dataflow Network

- Create a simple network
- See Debugging Aids
- See the TDF Debugger Visualizer



Demo

Network Builder Pattern

- You can build the network in the naïve way
 - Many LinkTo statements
- Or you can build the network using a build pattern
 - The Dataflow network is just a graph of objects
 - You can encapsulate the network building knowledge in a builder class
 - You can provide the building rule from the application logic and information file, metadata or use a DSL tool

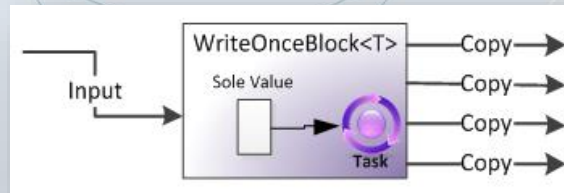
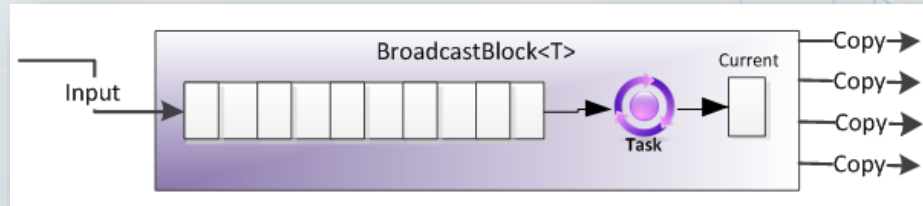
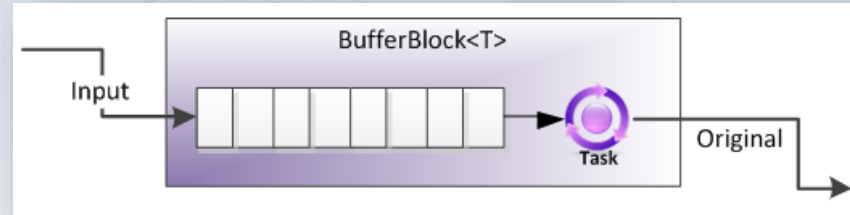
Data Block Options & Categories

- There are 9 public built-in data blocks in 3 categories:
 - **Buffering:** use to provide a buffer, collecting messages before dispatching them
 - Configured With [DataflowBlockOptions](#)
 - Capacity limit, Cancelation, Max messages per task, Task Scheduler
 - **Executing:** use to act upon a message
 - Configured With [ExecutionDataflowBlockOptions](#)
 - [DataflowBlockOptions](#) + Thread (Task) control
 - » Max degree of parallelism & Single producer constrained
 - **Grouping:** use to batch or join (or both) group of messages
 - Configured With [GroupingDataflowBlockOptions](#)
 - [DataflowBlockOptions](#) + Greedy consume messages, Max number of groups

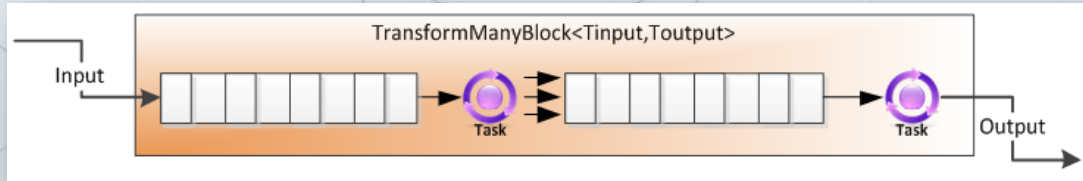
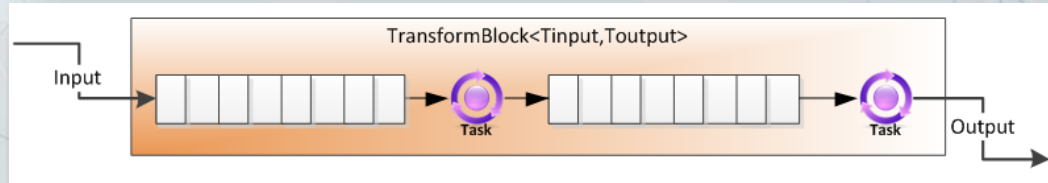
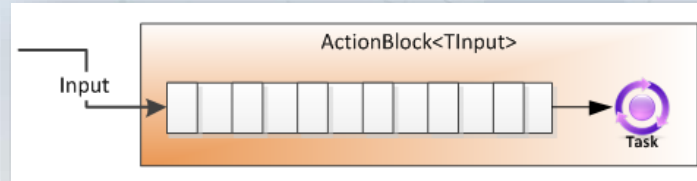
Predefine Blocks

Block	Type	Category	Description
<u>ActionBlock</u>	Target	Execution	Invokes an action for every message
<u>BatchBlock</u>	Propagator	Grouping	Batches input messages into an array
<u>BatchedJoinBlock</u>	Propagator	Grouping	Joins inputs to a batch of tuples
<u>BroadcastBlock</u>	Propagator	Buffering	Stores and distributes a copy of the last message
<u>BufferBlock</u>	Propagator	Buffering	Simple Buffer
<u>JoinBlock</u>	Propagator	Grouping	Form a tuple from a number of inputs
<u>TransformBlock</u>	Propagator	Execution	Transforms an input message to a different output
<u>TransformManyBlock</u>	Propagator	Execution	Transform an input message to number of outputs
<u>WriteOnceBlock</u>	Propagator	Buffering	Receiving and storing at most one element

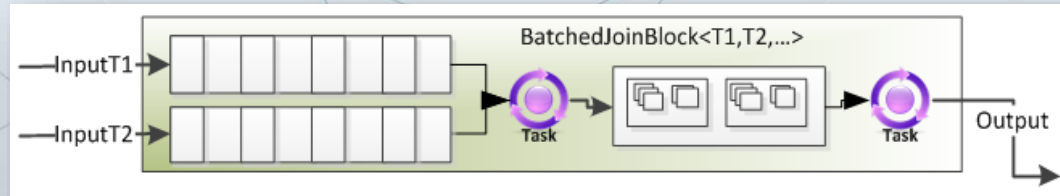
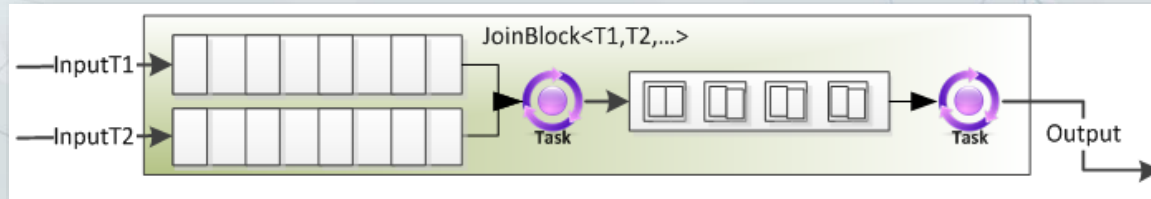
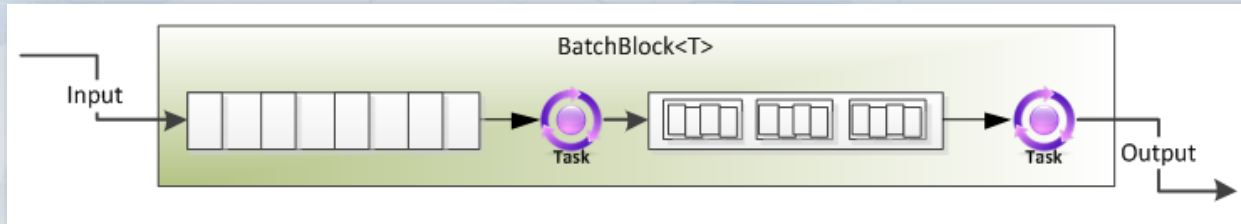
Blocks Internals - Pure Buffering



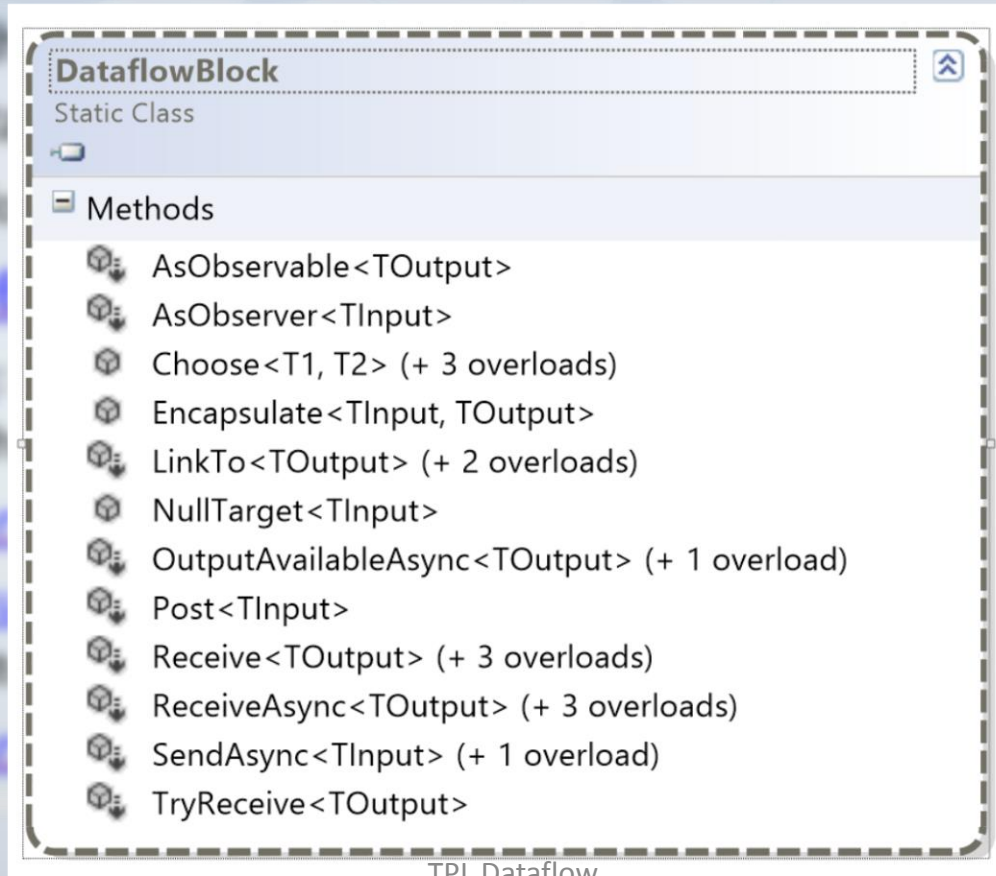
Block Internals – Execution



Block Internals - Grouping



The DataflowBlock Class



The screenshot shows the Visual Studio documentation for the `DataflowBlock` class. The class is identified as a 'Static Class'. Below this, a 'Methods' section lists the following members:

- `AsObservable<TOutput>`
- `AsObserver<TInput>`
- `Choose<T1, T2>` (+ 3 overloads)
- `Encapsulate<TInput, TOutput>`
- `LinkTo<TOutput>` (+ 2 overloads)
- `NullTarget<TInput>`
- `OutputAvailableAsync<TOutput>` (+ 1 overload)
- `Post<TInput>`
- `Receive<TOutput>` (+ 3 overloads)
- `ReceiveAsync<TOutput>` (+ 3 overloads)
- `SendAsync<TInput>` (+ 1 overload)
- `TryReceive<TOutput>`

The DataflowBlock Class

- DataflowBlock:
 - OutputAvailableAsync: informs of whether and when more output is available
 - LinkTo: Links the ISourceBlock to the specified ITargetBlock
 - Post: Posts an item to the ITargetBlock
 - Receive: Synchronously receives an item from ISourceBlock
 - TryReceive: Attempts to synchronously receive an item from the ISourceBlock
 - SendAsync: Asynchronously offers a message to ITargetBlock
 - ReceiveAsync: Asynchronously receives a value from ISourceBlock

Sending & Receiving Messages

- **Synchronous Blocking:**

```
block.SendAsync(m).Wait()  
m = block.Receive();
```

- **Synchronous non-blocking:**

```
while (!block.post(m)) ;  
while (!block.TryReceive(out m)) ;
```

- **Asynchronous:**

```
await block.SendAsync(m);  
m = await block.ReceiveAsync();
```

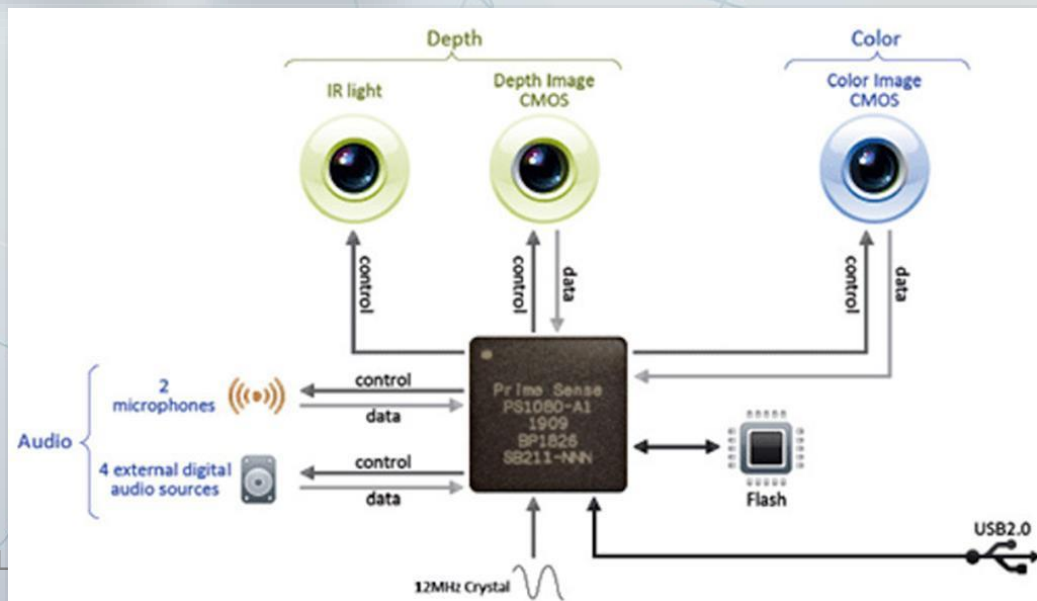
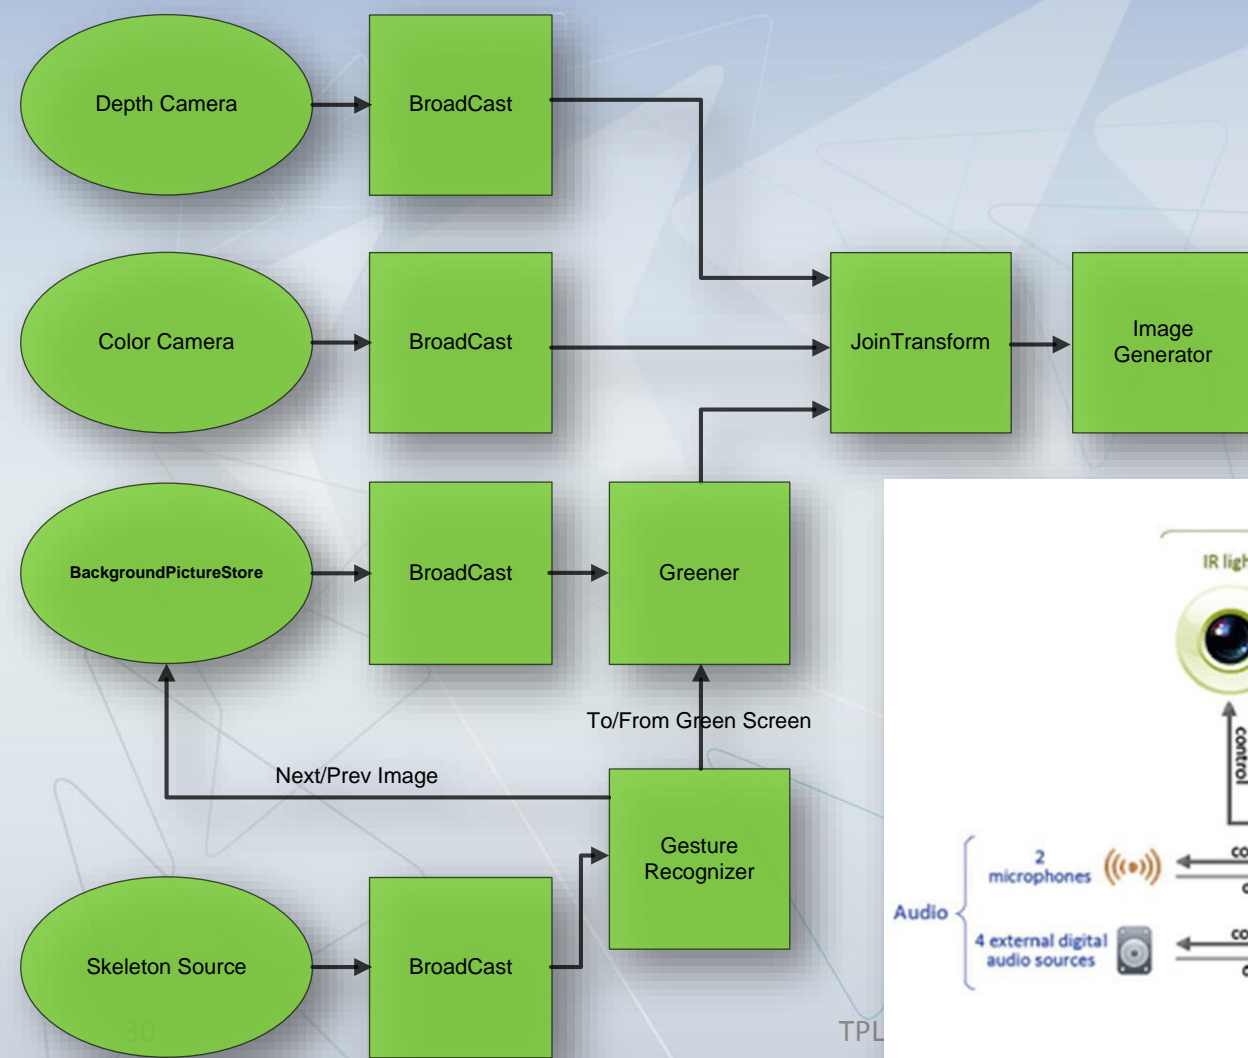
Concurrency Control

- Provide non-default scheduling schema such as priority
 - Specifying the Task Scheduler
- Concurrent execute multiple messages in execution blocks
 - Specifying the Degree of Parallelism
- Provide execution fairness among blocks
 - Specifying the Number of Messages per Task
- Cancel queued tasks
 - Provide TPL Cancellation token
- Remove unnecessary synchronization in execution block
 - Specify the a single producer drive the execution block
- Prevent deadlocks
 - Specify non-greedy behavior in grouping blocks

High Performance Low Latency Network



Demo



Handling Exceptions

- Block can become Faulted
 - Explicitly by calling to Fault
 - By calling a code that throws unhandled exception
 - By receiving a Fault from source block through Fault Propagation
 - By incorrect interface implementation
- Faulted block stop handling messages, clear its queues, fault its completion task and permanently declined messages
- Faulted Source block stop offering messages, propagates the faulted state (if it configured to do so) and unlink itself

Implementing Custom Blocks

- **Don't do it** if you don't have to
 - Prefer encapsulation with [DataflowBlock.Encapsulate](#)
 - Or embedding: build your block by implementing TDF interfaces and forwarding the messages to existing blocks
- Build the custom block the hard way by implementing TDF interfaces
- **It is not easy:**
 - Buffering, Messages Ownership, Synchronization, Block Completion, Non-Greediness, Reserve & Release

Summary

- We want our free lunch back
 - .NET TPL provide the needed parallel abstraction
 - TDF is a new way to architect, design & implement parallel execution following the Actor Model
- Dataflow network is built from dataflow blocks
 - You create the blocks and link them to form high performance low latency high throughput concurrent execution network
- You can customize existing blocks or even build your own

Q & A



Thank You